

Decode throughput, same GPU + model (RTX 5070 Ti, TinyLlama-1.1B)
viability, not supremacy: we lose on raw speed

